

Kickstarter Success Project

Course: CST338

Students: Jeff Burk, Aisha Lalli, Michelle Phung, Hugo Rodriguez

Instructor: Professor Memo

Introduction

Many entrepreneurs or small businesses turn to the Kickstarter platform as a means to fund their project. Significant time and effort can be put into creating a Kickstarter campaign, but if the funding goal is not reached then the project will not be funded, and could prove to have been a waste of resources. A number of factors can contribute to the success or failure of a Kickstarter campaign. Our goal is to create a model to predict the pledged dollar amount for a given campaign based on the dataset of over 300k completed campaigns, helping users decide whether to start a Kickstarter campaign or to pursue more traditional means of funding such as a bank loan or raising money from investors. Features used to predict the pledged amount include funding goal, main category, sub-category, length of campaign in days, and number of backers, and project state.

Dataset

Link to the dataset

<https://www.kaggle.com/kemical/kickstarter-projects> (<https://www.kaggle.com/kemical/kickstarter-projects>)

This dataset was collected from the Kickstarter platform. The dataset contains info on over 300K kickstarter projects from the company's launch in 2009 through 2017. Some features of the dataset include category, sub-category, project name, currency, country, funding goal, funding pledged, number of backers, project state, and more. The dataset includes 15 columns of which 7 are numeric values, and over 300k rows. The data was collected by a crowdfunding and data science enthusiast, Mickaël Mouillé, in 2018 and uploaded to Kaggle, an online community of data science enthusiasts. The dataset is currently available for download in Excel CSV format.

Module imports

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib import rcParams
import seaborn as sns
from scipy.stats import zscore
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsRegressor

from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split, cross_val_score
```

```
In [2]: # code in this cell from:
# https://stackoverflow.com/questions/27934885/how-to-hide-code-from-cells-in-ipython-notebook-visualized-with-nbviewer
from IPython.display import HTML

HTML("""<script>
code_show=true;
function code_toggle() {
  if (code_show){
    $('div.input').hide();
  } else {
    $('div.input').show();
  }
  code_show = !code_show
}
$(document).ready(code_toggle);
</script>
<form action="#">
  <input type="checkbox" /> Click here to display/hide the code.
</form>""")
```

Out[2]: ☐ Click here to display/hide the code.

Read the data and take a look at it

```
In [3]: infile = 'https://github.com/jb01010/cst383/raw/main/ks-projects-201801.zip'
df = pd.read_csv(infile)
```

Basic information of dataset

```
In [4]: df.describe()
```

```
Out[4]:
```

	ID	goal	pledged	backers	usd pledged	usd_pledged_real	usd_goal_real
count	3.786610e+05	3.786610e+05	3.786610e+05	378661.000000	3.748640e+05	3.786610e+05	3.786610e+05
mean	1.074731e+09	4.908079e+04	9.682979e+03	105.617476	7.036729e+03	9.058924e+03	4.545440e+04
std	6.190862e+08	1.183391e+06	9.563601e+04	907.185035	7.863975e+04	9.097334e+04	1.152950e+06
min	5.971000e+03	1.000000e-02	0.000000e+00	0.000000	0.000000e+00	0.000000e+00	1.000000e-02
25%	5.382635e+08	2.000000e+03	3.000000e+01	2.000000	1.698000e+01	3.100000e+01	2.000000e+03
50%	1.075276e+09	5.200000e+03	6.200000e+02	12.000000	3.947200e+02	6.243300e+02	5.500000e+03
75%	1.610149e+09	1.600000e+04	4.076000e+03	56.000000	3.034090e+03	4.050000e+03	1.550000e+04
max	2.147476e+09	1.000000e+08	2.033899e+07	219382.000000	2.033899e+07	2.033899e+07	1.663614e+08

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 378661 entries, 0 to 378660
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  --
0   ID                   378661 non-null  int64
1   name                 378657 non-null  object
2   category             378661 non-null  object
3   main_category        378661 non-null  object
4   currency             378661 non-null  object
```

View the first 5 rows of the dataset

In [6]: `df.head(5)`

Out[6]:

	ID	name	category	main_category	currency	deadline	goal	launched	pledged	state	backers	country	usd pledged	usd_pledged_real	usd_goal_real
0	1000002330	The Songs of Adelaide & Abullah	Poetry	Publishing	GBP	2015-10-09	1000.0	2015-08-11 12:12:28	0.0	failed	0	GB	0.0	0.0	1533.95
1	1000003930	Greeting From Earth: ZGAC Arts Capsule For ET	Narrative Film	Film & Video	USD	2017-11-01	30000.0	2017-09-02 04:43:57	2421.0	failed	15	US	100.0	2421.0	30000.00
2	1000004038	Where is Hank?	Narrative Film	Film & Video	USD	2013-02-26	45000.0	2013-01-12 00:20:50	220.0	failed	3	US	220.0	220.0	45000.00
3	1000007540	ToshiCapital Rekordz Needs Help to Complete Album	Music	Music	USD	2012-04-16	5000.0	2012-03-17 03:24:11	1.0	failed	1	US	1.0	1.0	5000.00
4	1000011046	Community Film Project: The Art of Neighborho...	Film & Video	Film & Video	USD	2015-08-29	19500.0	2015-07-04 08:35:03	1283.0	canceled	14	US	1283.0	1283.0	19500.00

Missing Data

Checking columns for nan values and calculating percentage

```
In [7]: df.isna().any()
percent_missing = df.isna().sum() * 100 / len(df)
pd.DataFrame({'column_name': df.columns, 'percent_missing': percent_missing})
```

Out[7]:

	column_name	percent_missing
	ID	0.000000
	name	0.001056
	category	0.000000
	main_category	0.000000
	currency	0.000000
	deadline	0.000000
	goal	0.000000
	launched	0.000000
	pledged	0.000000
	state	0.000000
	backers	0.000000
	country	0.000000
	usd pledged	1.002744
	usd_pledged_real	0.000000
	usd_goal_real	0.000000

A small percentage of the 'name' and 'usd_pledged' columns contain missing or NAN values These rows can be dropped

In [8]: `df.dropna(inplace=True)`

Preprocessing the data

Checking for duplicate project ID values, none found

```
In [9]: print("# of %s" % df[df.duplicated(subset='ID')].duplicated(subset='ID').sum())
0 Duplicate Project ID numbers
```

Dropping the project 'ID' column

In [10]: `df.drop(columns='ID', inplace=True)`

Checking the counts of campaign status

In [11]: `df.state.value_counts()`

Out[11]:

failed	197611
successful	133851
canceled	38757
live	2798
suspended	1843

Name: state, dtype: int64

Live, cancelled, suspended, and undefined campaigns can be dropped from the dataset

In [12]: `df = df[(df['state'] == 'successful') | (df['state'] == 'failed')]`

Displaying the updated value counts

In [13]: `df.state.value_counts()`

Out[13]:

failed	197611
successful	133851

Name: state, dtype: int64

Converting the launched and deadline dates to a timestamp format column for exploration and processing

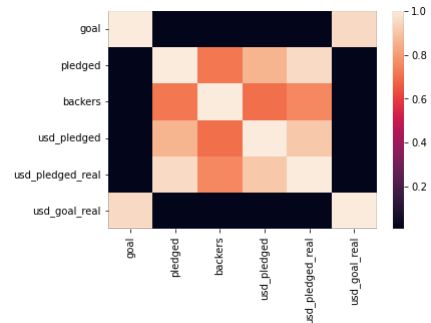
```
In [14]: df['launched'] = pd.to_datetime(df['launched'])
df['deadline'] = pd.to_datetime(df['deadline'])
```

Replacing column names that contain spaces for exploration and processing

```
In [15]: df.columns = [c.replace(' ', '_') for c in df.columns]
```

Finding data correlations of numeric columns with heatmap

```
In [16]: corr = df.corr()
sns.heatmap(corr, xticklabels=corr.columns, yticklabels=corr.columns):
```



The columns 'usd_goal_real' and 'usd_pledged_real' were added by the data collector to account for current currency conversion rates. These columns are highly correlated to 'goal' and 'usd_pledged' and can be dropped.

```
In [17]: df.drop(columns=['usd_goal_real', 'usd_pledged_real'], inplace=True)
```

Similarly 'usd_pledged' is highly correlated to 'pledged' and can be dropped


```
In [18]: df.drop(columns='usd_pledged', inplace=True)
```

The 'currency' column can also be dropped as we have the 'usd_pledged' values

```
In [19]: df.drop(columns='currency', inplace=True)
```

Of the remaining numeric columns, The 'goal' column appears uncorrelated to 'backers' or 'pledged'. The number of backers has about a 75% correlation to the amount pledged. Displaying updated heatmap.

```
In [20]: corr = df.corr()
sns.heatmap(corr):
```



Additional features are kept for data exploration but will be dropped to prepare for machine learning. Displaying remaining features and info.

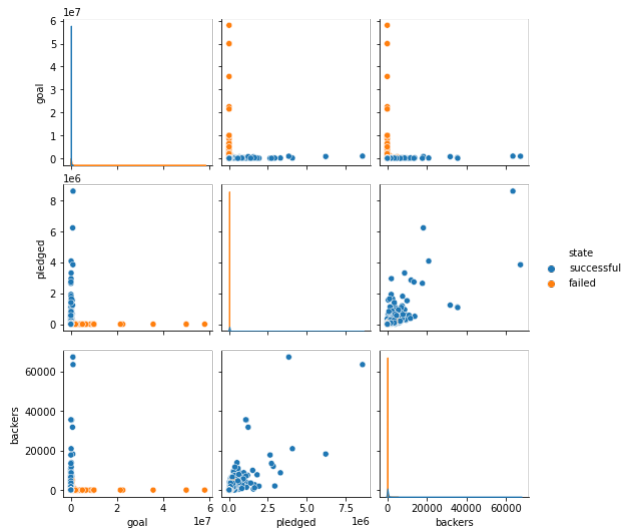
```
In [21]: df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 331462 entries, 0 to 378660
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   name             331462 non-null object
1   category         331462 non-null object
2   main_category    331462 non-null object
3   deadline         331462 non-null datetime64[ns]
4   goal             331462 non-null float64
5   launched         331462 non-null datetime64[ns]
6   pledged          331462 non-null float64
7   state            331462 non-null object
8   backers          331462 non-null int64
9   country          331462 non-null object
dtypes: datetime64[ns](2), float64(2), int64(1), object(5)
memory usage: 27.8+ MB
```

Exploring the data

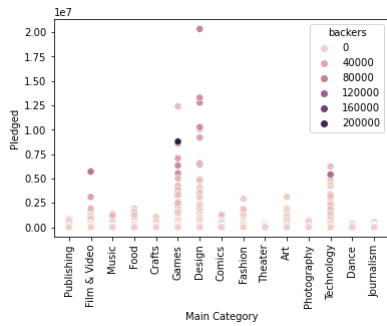
Creating pairplot to view/confirm data correlations. The columns 'backers' and 'pledged' appear to be correlated

```
In [22]: # Can be slow, creating a random sampling to speed up
df_mini = df.sample(n=25000)
sns.pairplot(df_mini, hue="state", height=2.5)
```




Creating a scatterplot of the amount pledged for each main category

```
In [23]: sns.scatterplot(x='main_category', y='pledged', hue='backers', data=df, s=50);
plt.xticks(rotation=90);
plt.xlabel('Main Category');
```




Creating a boxplot of amounts pledged by year

```
In [24]: # Log transform 'pledged' to distribute values
np.seterr(divide='ignore');
plt.figure(figsize=(13,4));
sns.boxplot(x=df.launches.dt.year, y=np.log(df.pledged));
plt.xlabel('Year launched');
plt.ylabel('Amount pledged');
plt.title('Amount pledged by year');
```



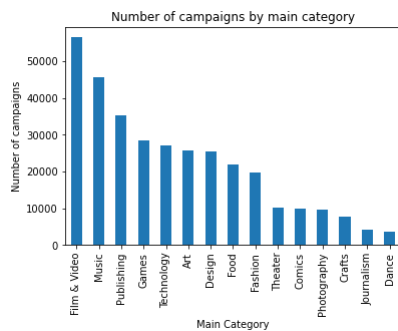
Creating barplot for number of campaigns by country

```
In [25]: df['country'].value_counts()
df['country'].value_counts().plot.bar();
plt.xlabel('Country');
plt.ylabel('Number of campaigns');
plt.title('Number of campaigns by country');
```



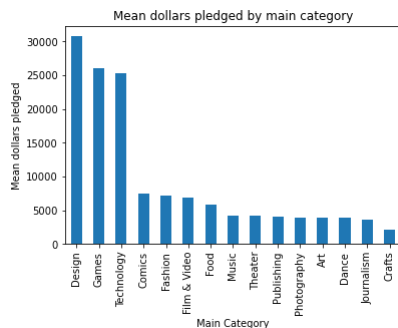
Creating barplot for number of campaigns by main category

```
In [26]: df['main_category'].value_counts().plot(kind='bar', title='Number of Campaigns by Main Category');
plt.xlabel('Main Category');
plt.ylabel('Number of campaigns');
plt.title('Number of campaigns by main category');
```



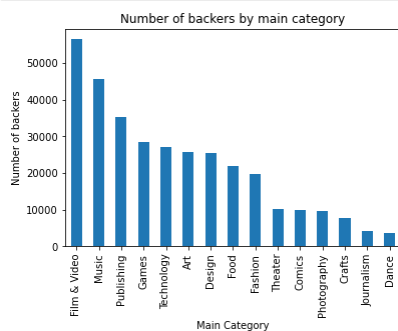
Creating barplot for mean dollars pledged by main category

```
In [27]: df.groupby('main_category').mean()['pledged'].sort_values(ascending=False).plot(kind='bar');
plt.xlabel('Main Category');
plt.ylabel('Mean dollars pledged');
plt.title('Mean dollars pledged by main category');
```



Creating barplot for number of backers by main category

```
In [28]: df.groupby('main_category')['backers'].count().sort_values(ascending=False).plot(kind='bar');
plt.xlabel('Main Category');
plt.ylabel('Number of backers');
plt.title('Number of backers by main category');
```



Display how many and list categories within main categories


```
In [29]: df.groupby('main_category')['category'].value_counts()
```

Out[29]:

```
main_category  category
Art            Art      7486
              Painting  3034
              Illustration 2890
              Public Art  2850
              Mixed Media 2513
              ...
Theater        Festivals 512
              Experimental 342
              Immersive   297
              Soaces      194
```

Number of backers by sub category, under 50k backers for legibility

```
In [30]: plt.figure(figsize=(25,5))
df2 = df[df['backers'] < 50000]
sns.scatterplot(x='category', y='backers', hue='main_category', data=df2);
plt.xlabel('Sub Category');
plt.xticks(rotation=90);
plt.ylabel('Number of backers');
plt.title('Number of backers by sub category, under 50k backers');
plt.legend(loc='best');
```



Preparing data for machine learning

```
In [31]: np.random.seed(78)
```

Converting categorical features to numeric. The best test RMSE was about 15% lower when converting the categorical columns manually to percentages versus using the `get_dummies()` method, with 6 predictors.

```
In [32]: # Get campaign Length in days
df['days'] = (df.deadline - df.launched).astype('timedelta64[D]').astype(int)

# Create numeric representation of 'main_category' by calculating mean percent of pledged amounts
def func(x):
    x['m_cat_perc'] = x['pledged'].mean() / x['pledged'].sum().sum()
    return x
df = df.groupby('main_category').apply(func)

# Create numeric representation of 'category' by calculating mean percent of pledged amounts
def func(x):
    x['s_cat_perc'] = x['pledged'].mean() / x['pledged'].sum().sum()
    return x
df = df.groupby('category').apply(func)

df['state'] = (df['state'] == 'successful').values.astype(int)

# predictors = ['goal', 'backers', 'days', 'category', 'main_category', 'state']
# target = 'pledged'
# X = df[predictors]
# y = df[target]
```

```
In [33]: # select predictor and target variables to be used with regression model, dropping other features
predictors = ['goal', 'backers', 'days', 'm_cat_perc', 's_cat_perc', 'state']
# predictors = ['goal', 'backers', 'days', 'category', 'main_category', 'state', 'country']
target = 'pledged'
X = df[predictors].values
y = df[target].values
```

Getting a random sample of the data as knn can be slow

```
In [34]: indexes = np.random.choice(y.size, size=15000)
X_mini = X[indexes]
y_mini = y[indexes]

# X_mini = X_mini[(np.abs(zscore(X_mini))) < 3].all()]
# y_mini = y_mini[(np.abs(zscore(y_mini))) < 3].all()]
```

Scaling the data

```
In [35]: # Split into training and test sets and scale
scaler = StandardScaler()

# unscaled version, scaling is only used on predictor variables
X_train_raw, X_test_raw, y_train, y_test = train_test_split(X_mini, y_mini, test_size=0.25, random_state=1)

# scaled version
X_train = scaler.fit_transform(X_train_raw)
X_test = scaler.transform(X_test_raw)

#X_train = X_train[(X_train < 3) | (X_train > -3)]
#X_test = X_test[(X_test < 3) | (X_test > -3)]
```

Data sanity check

```
In [36]: print(X_train.shape)
print(X_train[:3])
(11250, 6)
[[-0.0559271 -0.17370337 -0.30750036 -0.41657818 -0.25264235 -0.81785809]
 [-0.01585716 -0.1703651  2.03424273 -0.15026109 -0.37779915 -0.81785809]
 [-0.05609406 -0.14032067 -0.30750036 -0.56901524 -0.2625667  1.22270601]]
```

Baseline performance

Getting the baseline "blind" prediction, average value of the target variable

```
In [37]: def rmse(predicted, actual):
return np.sqrt(((predicted - actual)**2).mean())
```

```
In [38]: actual = y_test
predicted = y_train.mean()
print("Test RMSE baseline: {}".format(rmse(predicted, y_test)))
test, rmse baseline: 66072.9
```

KNeighborsRegressor - Performance with default hyperparameters

```
In [39]: knn = KNeighborsRegressor(algorithm='brute')
knn.fit(X_train, y_train)
predicted = knn.predict(X_test)
print("Test RMSE, default hyperparameters: {}".format(rmse(predicted, y_test)))
test RMSE, default hyperparameters: 52169.5
```

Testing with different values of K

```
In [40]: def get_train_test_rmse(knn, X_train, X_test, y_train, y_test):
knn.fit(X_train, y_train)
predictions = knn.predict(X_test)
train_rmse = rmse(y_train, knn.predict(X_train))
test_rmse = rmse(y_test, predictions)
return (train_rmse, test_rmse)

In [41]: n = 30
test_rmse = []
train_rmse = []
ks = np.arange(1, n+1, 2)
for k in ks:
    print(k, ' ', end='')
    knn = KNeighborsRegressor(n_neighbors=k)
    rmse_tr, rmse_te = get_train_test_rmse(knn, X_train, X_test, y_train, y_test)
    train_rmse.append(rmse_tr)
    test_rmse.append(rmse_te)
print('\ndone!')
```

1 3 5 7 9 11 13 15 17 19 21 23 25 27 29 done

Print test RMSE results as a function of k, for all odd value k up to 29

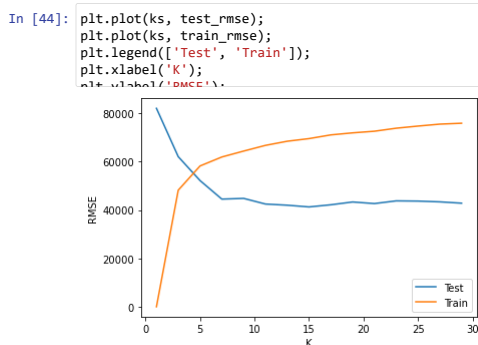
```
In [42]: print(test_rmse)
[81997.89121360613, 62029.24584675199, 52169.54031329579, 44517.03310098852, 44812.34500545312, 42485.43682704381, 42003.24968350373, 41291.69995824102, 42180.3238300688
7, 43331.88123336618, 42677.29482436208, 43791.71227635761, 43700.95751088592, 43385.6035475931, 42837.385333659724]
```

Print the best k and the test RMSE associated with it

```
In [43]: def get_best(ks, rmse):
best = rmse[0]
bestk = ks[0]
for i in range(0, len(ks)):
    if (rmse[i] < best):
        best = rmse[i]
        bestk = ks[i]
    #print(i)
    #print(rmse[i])
return int(bestk), round(best,1)

best_k, best_rmse = get_best(ks, test_rmse)
print('best k = {}, best test RMSE: {}'.format(best_k, best_rmse))
best k = 15, best test RMSE: 41291.7
```

Plot the test and training RMSE as a function of k, for all odd value k up to 29



The best RMSE while testing hyperparameter values of k on the KNeighborsRegressor model was 41291.7

LinearRegressor model

```
In [45]: # select predictor and target variables to be used with model, dropping other features
predictors = ['goal', 'backers', 'days', 'm_cat_perc', 's_cat_perc', 'state']
target = 'pledged'

X = df[predictors].values
y = df[target].values

reg = LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
reg.fit(X, y)
```

Out[45]: LinearRegression()

View the coefficients and intercept of the LinearRegression model

```
In [46]: print(reg.intercept_)
print(reg.coef_)
# print('intercept: ' + str(round(reg.intercept_, 2)))
# print('coefficients: ')
-2396.2822501874325
[ 4.31160035e-04  7.52031247e+01  8.28152528e+01 -1.10918978e+07
  4.57663281e+05  4.34176191e+03]
```


Plotting the actual values versus the predicted values

```
In [47]: def plot_actual_predicted(actual, predicted, title):
sns.scatterplot(x=actual[:], y=predicted[:], s=40)

x = predicted[:].min(), actual[:].max()
y = predicted[:].min(), actual[:].max()


plt.plot(x, y, color='black', linestyle='dashed');
plt.title(title);
plt.xlabel('actual');
plt.ylabel('predicted');
```

`plot_actual_predicted(y_train, reg2.predict(X_train), 'Predicted by actual (training data)')`




Many predictions are off the line. Fitting another model after splitting into training and test data.

```
In [48]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=1)
reg2 = LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
reg2.fit(X_train, y_train)
plot_actual_predicted(y_train, reg2.predict(X_train), 'Predicted by actual (training data)')
```



```
In [49]: plot_actual_predicted(y_test, reg2.predict(X_test), 'Predicted by actual (test data)')
```



The plots of test and training data appear similar, checking RMSE and r-squared values

```
In [50]: RMSE = np.sqrt(((reg2.predict(X_test) - y_test)**2).mean())
print(round(RMSE, 2))
print('r-squared value of reg2: f: {f1}' format(reg2.score(X_train, y_train)))
47526.27
r-squared value of reg2: 0.5069
```

The new model increases the RMSE by about 15%.


```
In [51]: predictors = ['goal', 'backers', 'days', 'm_cat_perc', 's_cat_perc', 'state']

X = np.array(df[predictors])
X_s = zscore(X)
y = df['raised']
```

Let's check the r-squared value of the scaled data

```
In [52]: X_train, X_test, y_train, y_test = train_test_split(X_s, y, test_size=0.25, random_state=1)

reg3 = LinearRegression()
reg3.fit(X_train, y_train)

print('r-squared value of reg2: {}'.format(reg3.score(X_train, y_train)))
r-squared value of reg3: 0.5069
```

Scaling the data has not affected the results, let's try adding some Polynomial Features

```
In [53]: pf = PolynomialFeatures(degree=2)
pf.fit(X)
X_poly = pf.transform(X)
```

Checking the shape

```
In [54]: X_poly.shape
Out[54]: (331462, 28)
```

Creating a model using all of the new features

```
In [55]: X_train, X_test, y_train, y_test = train_test_split(X_poly, y, test_size=0.25, random_state=1)

reg4 = LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
reg4.fit(X_train, y_train)

predict = reg4.predict(X_test)
RMSE = np.sqrt(((predict - y_test)**2).mean())
print('RMSE: {}'.format(RMSE))
RMSE: 46128.66
```

Check the RMSE of the first feature

```
In [56]: X_0 = X_train[:, [0]]

# compute the negated mse scores with 5-fold cross validation
scores = cross_val_score(LinearRegression(), X_0, y_train, scoring='neg_mean_squared_error', cv=5)
RMSE = np.sqrt(-scores.mean())

print('RMSE: {}'.format(RMSE))
RMSE: 109325.81
```

Find feature with the lowest RMSE

```
In [57]: rmse_min = []

for i in range(X_train.shape[1]):
    X_0 = X_train[:, [i]]
    scores = cross_val_score(LinearRegression(), X_0, y_train, scoring='neg_mean_squared_error', cv=5)
    rmse_min.append(np.sqrt(-scores.mean()))
i_min = np.where(rmse_min == np.min(rmse_min))[0][0]

rmse_min = np.min(rmse_min)
print('best feature: {}, best RMSE: {}'.format(i_min, rmse_min))
best feature: x1, best RMSE: 78607.68
```

Now let's find the best set of 8 features using the greedy method of forward feature selection.

```
In [58]: remaining = list(range(X_train.shape[1]))
selected = []

n = 8
while len(selected) < n:
    rmse_min = 1e7
    for i in remaining:
        selected.append(i)
        X = X_train[:, selected]
        scores = cross_val_score(LinearRegression(), X, y_train, scoring='neg_mean_squared_error', cv=5)
        selected.remove(i)
        if (np.sqrt(-scores.mean()) < rmse_min):
            rmse_min = np.sqrt(-scores.mean())
            i_min = i
    remaining.remove(i_min)
    selected.append(i_min)
    print('num features: {}, rmse: {}'.format(len(selected), rmse_min))

num features: 1; rmse: 78607.68
num features: 2; rmse: 74663.02
num features: 3; rmse: 72597.77
num features: 4; rmse: 72419.93
num features: 5; rmse: 72317.89
num features: 6; rmse: 72211.72
num features: 7; rmse: 72205.64
num features: 8; rmse: 72204.96
```

Shows little improvement after 7 features. Check RMSE of 8 feature set

```
In [59]: X = X_poly[:,selected]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=1)

reg7 = LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
reg7.fit(X_train, y_train)

predict = reg7.predict(X_test)
RMSE = np.sqrt(((predict - y_test)**2).mean())
print('test RMSE with 9 features: {}'.format(RMSE))
test RMSE with 8 features: 46397.6
```

Our model using all features has only slightly better RMSE value than the 8 feature set

DecisionTreeRegressor model

```
In [60]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=1)

reg8 = DecisionTreeRegressor(max_depth=5)
reg8.fit(X_train, y_train)

y_predict = reg8.predict(X_test)
errors = y_test - y_predict
rmse = np.sqrt((errors**2).mean())
print('rmse: {}'.format(rmse))
rmse: 44574.83
```

Conclusion

The best RMSE while testing hyperparameter values of k on the KNeighborsRegressor model was 41291.7 with a k value of 15. The best RMSE while testing the LinearRegressor model was 46128.66 using all polynomial features and 46397.6 using the top 8 features. The best RMSE while testing max depth values of the DecisionTreeRegressor model was 44574.83 with a max depth of 5.

In []: